

Lab 5

It's all about the Processors

Introduction

In this lab, we are going to use almost everything we have learned and created thus far in order to make the most recognizable design in all of ECE; a processor. Our design is a general purpose processor with some application specific instructions for video and communications. In this way, it is similar to a simplified Application-Specific Instruction-Set Processor (ASIP).

In order to implement it, a fully detailed Instruction Set Architecture (ISA) has been created that will be the core of the design. By modifying previous components to fit our needs and leveraging the memory IP available in the Vivado tool, an entire computer system will be created with relatively little effort on our part. Once the design has been created, it will be tested by running programs that were written in assembly and passed through a custom assembler as well as running it through a full system simulator.

Prelab - What Kind of Chip you got in there?

Background

Before we can go about making our processor, we have to understand how it works. To do so, we must examine the ISA that has been created. This will tell us the available instructions, memory organization, register file structure, and instruction formats.

In general, a processor consists of a group of “dumb” components, like memories, general-purpose registers, multiplexers, and an ALU, and a “smart” controller. By using a Finite State Machine as the controller, we can put those components to work to accomplish meaningful tasks.

For the bulk of this lab, we will be creating and modifying these components in order to prepare them for our top level design. Then we will create the most complicated part; the controller.

Tasks

- Read through the ISA specification attached at the end of the manual.
- Watch this video on the Vivado IP Integrator Tool:
<https://www.youtube.com/watch?v=ZXiygbhmZoE>
- Modify your entity “my_alu” from Lab 2 to have a width of 16 bits and the following opcodes given in table 5.1. Make it synchronous with a clock enable input.
- Modify your entity “pixel_pusher” from Lab 4 to output a 64x64 resolution image from hcount = [0, 63] and vcount = [0, 63] using a 12-bit “addr”. It will read in a 16-bit pixel instead of an 8-bit one with the following distribution: [R,G,B] = [5,6,5]

Table 5.1: Modified ALU Instructions

opcode	function	opcode	function
x“0”	$A + B$	x“8”	$A \text{ and } B$
x“1”	$A - B$	x“9”	$A \text{ or } B$
x“2”	$A + 1$	x“A”	$A \text{ xor } B$
x“3”	$A - 1$	x“B”	$A < B$ (signed) (as bit 0 of output)
x“4”	$0 - A$	x“C”	$A > B$ (signed) (as bit 0 of output)
x“5”	$A << 1$ (shift left logical)	x“D”	$A = B$ (as bit 0 of output)
x“6”	$A >> 1$ (shift right logical)	x“E”	$A < B$ (as bit 0 of output)
x“7”	$A >>> 1$ (shift right arithmetic)	x“F”	$A > B$ (as bit 0 of output)

1 Wanna be Hackers? Code Crackers? Slackers?

1.1 Background

In order to use our processor, we're going to need some programs to run on it. In order to get those, we're going to write them ourselves. However, writing binary instructions is really annoying and time-intensive. Instead, two versions of an assembler have been prepared that will convert raw text assembly programs into binary files to use with our design. One was written using MATLAB, while the other was written using Python.

For our test programs, we will do simple tasks that exercise the IO interfaces available in the architecture, the UART and the VGA display. By using these interfaces with some simple programs, we can easily prove the basic functionality of the design.

In order to understand the format of the assembly language, please look at the provided text files. Please note that the syntax varies slightly depending on which assembler you choose to use. Look at the comments in the assembler source for a description of the allowed formats.

1.2 Tasks

- Write an assembly program that prints the string "Hello-World" over the UART, one byte at a time by reading characters from a string declared in the data segment and sending them until it encounters a null character. After it finishes, it should continuously loop through reading a character from the UART into a register and writing that register value to all memory locations in the video memory. Assemble it and store the text and data COE files for later use. Simulate it to verify code functionality using the provided simulator.

2 Every File Inspected, no Viruses Detected

2.1 Background

In order for our processor to have data to work with, we need memories. For simplicity, our system uses a pure Harvard Architecture with separate instruction and data memories. In this way, our executable code is stored in a single port read-only memory that is read from by the controller. Our data memory is stored in a single port ram memory that allows read/write access. Additionally, we implement our Register File as a dual port memory in order to allow double reads for R-type instructions. Finally, our VGA controller reads data from a dual port ram known as a framebuffer (which can be written and read to by the CPU) in order to output the contents of the framebuffer to the screen.

Some of these memories (the instruction and data memories) will be auto generated, in a similar method to the one used in Lab 4, in order to "program" them in advance with the outputs of our assembler. For the others, we are going to manually create them ourselves using simple VHDL that infers memory blocks.

In order to infer memory, we start by modeling our memory as an array of `std_logic_vectors`, similar to how a ROM was inferred in Lab 3 for your NetID. By describing the behavior of operations on indexes in that signal via some clever indexing and type conversions (think Lab 3), we can infer memories that the tool will then replace with the appropriate hardware primitives.

2.2 Tasks

- Create an entity called “regs” where you infer a true dual port (both ports can independently either read or write to any location) memory consisting of 32 16-bit words (64 Bytes). It should write synchronously and read asynchronously. See listing 5.1 for a black box description, and listing 5.2 for behavior.
- Create an entity called “framebuffer” where you infer a dual port memory consisting of 4096 16-bit words (8 KiB). See listing 5.3 for the black box interface / entity declaration. *[Notice how we have two different enables going into this device. The CPU side and the video display side are on two different clock domains, operating independently of each other. Because one side is only reading we can safely do this on a dual port memory. Otherwise we would have to navigate access with a controller and some FIFOs]*

Listing 5.1: Register Entity

```
entity regs is port (
  clk, en, rst      : in std_logic;
  id1, id2          : in std_logic_vector(4 downto 0);      -- Addresses
  wr_en1, wr_en2    : in std_logic;
  din1, din2        : in std_logic_vector(15 downto 0);
  dout1, dout2      : out std_logic_vector(15 downto 0)
);
end regs;
```

Listing 5.2: Register Description

```
# sudo code for register behavior

dout1 = registers(id1)
dout2 = registers(id2)

if on rising edge of clock:
  registers(0) = 0

  if reset is 1:
    registers(all) = 0

  else if enable is 1:

    if wr_en1:
      registers(id1) = din1

    if wr_en2:
      registers(id2) = din2
```

Listing 5.3: Framebuffer Entity

```

entity framebuffer is
port (
    clk1, en1, en2, ld    : in std_logic;
    addr1, addr2          : in std_logic_vector(11 downto 0);
    wr_en1                : in std_logic;
    din1                  : in std_logic_vector(15 downto 0);
    dout1, dout2          : out std_logic_vector(15 downto 0)
);
end framebuffer;

```

See the following VHDL design for framebuffer reference:

Listing 5.4: Dual Port RAM Infer Sample

```

-- a parameterized, inferable, tdp, dual-clock bram in vhdl

library ieee;
    use ieee.std_logic_1164.all;
    use ieee.numeric_std.all;

entity bram_tdp is
generic (
    DATA    : integer := 72;
    ADDR     : integer := 10
);
port (
    -- port a
    a_clk    : in std_logic;
    a_wr     : in std_logic;
    a_addr   : in std_logic(ADDR-1 downto 0);
    a_din    : in std_logic(DATA-1 downto 0);
    a_dout   : out std_logic(DATA-1 downto 0);

    -- port b
    b_clk    : in std_logic;
    b_wr     : in std_logic;
    b_addr   : in std_logic(ADDR-1 downto 0);
    b_din    : in std_logic(DATA-1 downto 0);
    b_dout   : out std_logic(DATA-1 downto 0)
);
end bram_tdp;

architecture rtl of bram_tdp is

    -- memory
    type mem_type is array (0 to (2**ADDR)-1) of std_logic_vector(DATA-1 downto
        ↪ 0);
    signal mem : mem_type;

begin

    -- port a
    process(a_clk)

```

```

begin
    if rising_edge(a_clk) then
        if (a_wr = '1') then
            mem(to_integer(unsigned(a_addr))) <= a_din;
        end if;
        a_dout <= mem(to_integer(unsigned(a_addr)));
    end if;
end process;

-- port b
process(b_clk)
begin
    if rising_edge(b_clk) then
        if (b_wr = '1') then
            mem(to_integer(unsigned(b_addr))) <= b_din;
        end if;
        b_dout <= mem(to_integer(unsigned(b_addr)));
    end if;
end process;

end rtl;

```

Hint: Our goal with framebuffer is to infer BRAM. In order to fully reset it (i.e. zero out the memory), you will need to get creative. Attempting to reset all memory locations simultaneously will cause the tool to synthesize framebuffer as registers, which is not what we want.

3 Where'd you get your CPU?

3.1 Background

We have finally reached it, the big kahuna. In this section we are going to make the brains of the operation, the controller. Everything ends up becoming a state in our machine and adding custom instructions is as simple as adding more states to the machine. This makes the controller mostly mindless grunt work, since the FSM states themselves are relatively simple. Our FSM models a multi-cycle CPU with an average CPI of roughly 5 clock cycles.

3.2 Tasks

- Create the FSM for the control unit. See listing (5.5) for entity declaration. In order to facilitate the development of the controller, a simplified state diagram and state behavior table have been given in figure (5.2). It will need more states in some areas due to latency with memories and the ALU (see notes below).
- Simulate your controller for a simple arithmetic instruction by asserting the instruction “add \$r3 \$r4 \$r5” (x“00C85000”). Verify that it goes through the proper states and asserts the correct control signals.

Listing 5.5: Controls Entity Description

```
entity controls is
port (
  -- Timing Signals
  clk, en, rst : in std_logic;

  -- Register File IO
  rID1, rID2      : out std_logic_vector(4 downto 0);
  wr_enR1, wr_enR2 : out std_logic;
  regrD1, regrD2   : in std_logic_vector(15 downto 0);
  regwD1, regwD2   : out std_logic_vector(15 downto 0);

  -- Framebuffer IO
  fbRST      : out std_logic;
  fbAddr1    : out std_logic_vector(11 downto 0);
  fbDin1     : in std_logic_vector(15 downto 0);
  fbDout1    : out std_logic_vector(15 downto 0);

  -- Instruction Memory IO
  irAddr : out std_logic_vector(13 downto 0);
  irWord : in std_logic_vector(31 downto 0);

  -- Data Memory IO
  dAddr  : out std_logic_vector(14 downto 0);
  d_wr_en : out std_logic;
  dOut   : out std_logic_vector(15 downto 0);
  dIn    : in std_logic_vector(15 downto 0);

  -- ALU IO
  aluA, aluB : out std_logic_vector(15 downto 0);
  aluOp      : out std_logic_vector(3 downto 0);
  aluResult  : in std_logic_vector(15 downto 0);

  -- UART IO
  ready, newChar : in std_logic;
  send          : out std_logic;
  charRec       : in std_logic_vector(7 downto 0);
  charSend      : out std_logic_vector(7 downto 0)
);
end controls
```

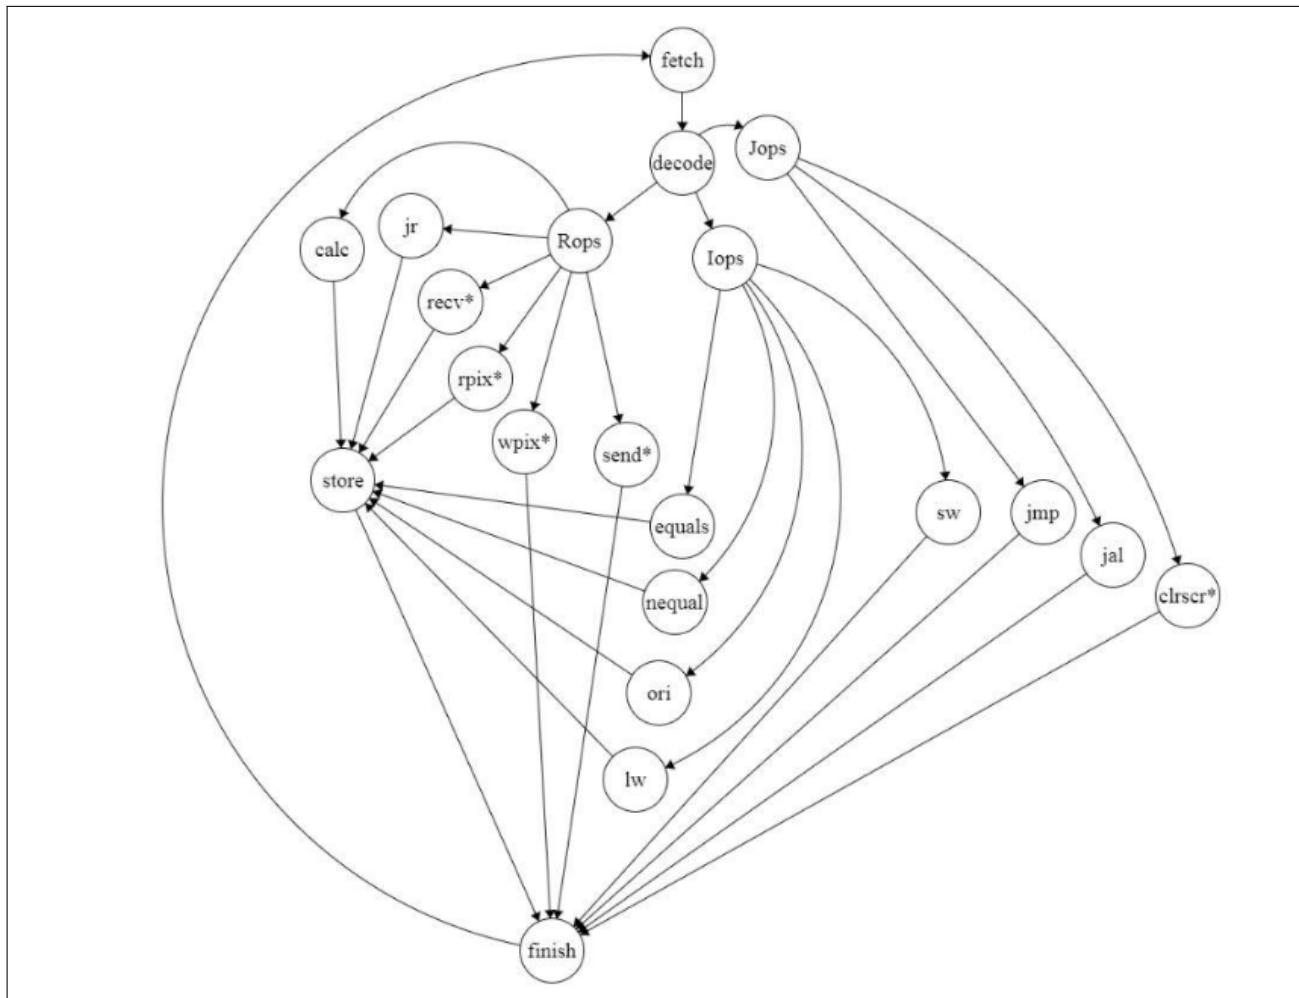


Figure 5.1: Simplified State Diagram for “controls” (asterick denotes ASIP instruction)

NOTE 1: Reading from irMem and dMem has a 1 clock cycle latency. Your results will not be available until the next rising edge after you assert your control signals. This 1 cycle latency also exists for reading the result from your ALU

NOTE 2: Any control signals asserted should be immediately de-asserted on the next state in your state machine

Table 5.2: “Controls” State Transition and Behavior Table

Current State	Actions	Next State
fetch	get current pc from reg into signal	decode
decode	store irMem[pc_signal] into a signal store pc_signal+1 into register 1	Rops if opcode top bits are 00 or 01 else Iops if 10 else Jops
Rops	Break up instruction into arguments. Use arguments to fetch register contents for reg2 and reg3 into signals	jr if opcode is 01101 else recv if 01100 else rpix if 01111 else wpix if 01110 else send if 01011 else calc
Iops	Break up instruction into arguments. Use arguments to fetch register contents for reg2 into signal	equals if opcode bottom 3 bits are 000 else nequal if 001 else ori if 010 else lw if 011 else sw
Jops	Break up instruction into arguments	jmp if opcode is 11000 else jal if 11001 else clrscr
calc	Apply the register operands and the cor- rect opcode to the ALU and store the result into an alu result signal	store
store	Store the alu result signal into the ap- propriate register given by argument reg1	finish
jr	Read the register value specified and store it in alu result signal	store
recv (asip)	Store charRec into alu result signal	recv if newChar is 0 else store
rpix (asip)	Read the framebuffer memory at the ad- dress of the value in reg2 and store it in alu result signal	store
wpix (asip)	Store the value read from reg2 into framebuffer[reg1]	finish
send (asip)	Make send 1 and assign the value read from reg1 to charSend	If ready is 1 finish, else send
equals	If values equal set alu signal to immedi- ate and set reg1 signal to pc id	store
nequal	If values not equal set alu signal to im- mediate and set reg1 signal to pc id	store

Current State	Actions	Next State
ori	Store the result of the immediate bitwise ORed with the value from reg2 into the alu signal	store
lw	Set the value of the alu signal to the value in dmem[reg2+imm]	store
sw	Store the value of reg1 into dmem[reg2+imm]	finish
jmp	Set the value of the pc register to the immediate	finish
jal	Set the value of the ra register to the value of the pc register and set the value of the pc register to the immediate	finish
clrscr (asip)	Set fbRST to 1	finish
finish	De-assert any required control signals	fetch

4 This isn't even my final form

4.1 Background

In order to implement this design on the board, which is clearly gargantuan, we are going to take advantage of a graphical entry tool called the IP Integrator. In order to start working in the IP Integrator, click “Create Block Design” under the IP Integrator flow tab. Specify some arbitrary name for you design and click OK. This will bring you into the integrator window. From here, drop in your designs by right clicking inside the main window and selecting “Add Module” double click the module you wish to add and it will be dropped into the block diagram view. Ignore all connection automation prompts.

Once all your modules have been dropped in, you can connect them by clicking on a port and dragging the connection to another port. In order to make ports go to your top level IO, right click and select “Make External”. Then rename the external contact to match the name in your XDC file by right clicking the port, selecting “External Port Properties”, and renaming the port name.

After connecting all of your modules, you will need to add ports that are tied to high impedance (Our CTS and RTS signals for the UART PMOD). In order to do so, go to the ports list in the left hand helper window. Right click the “ports” section and select “Create port”. Set the name of your top level port and the direction. Then click OK. That's it, it stays unconnected to be tied to high impedance via a tri-state buffer.

Now click the checkmark icon in order to validate your block design and check for any errors. Once that

is done the design is finished.

Finally, go back to the “Project Manager” flow tab. Right click your block diagram file and select “Create HDL Wrapper”. The tool will generate warnings about width mismatches for our memories and our controller. This is fine. It will only connect the lower bits together, which is what we want. Then set the VHDL file it creates as your top level design. You are now ready to proceed to synthesis.

4.2 Tasks

- Create an appropriately modified XDC file.
- Implement your design in a top level according to the simplified block diagram given below in order to test it on the board using the IP Integrator Tool described in the background section above. Please include the block diagram image you generate in your final lab report.
- Create a single port rom, called “irMem” (Our instruction memory), using the Xilinx Block Memory IP in the same manner as described in Lab 4 Part 2. Set the width to 32 bits and the depth to 16384 (a total of 64 KiB). [Notice how this memory is not large enough to have a 16 bit address space since we are addressing by double word rather than byte. Our processor has a “virtual” address space of 16 bits for the instruction memory but our physical system only has a 14 bit instruction address space due to memory limitations. In our top level design, we will ignore the upper 2 bits of our address when addressing this memory]. Disable any output registering and set the enables to always enabled. In the “Other Options” tab, set it to initialize using the text coe file generated in Part 1.
- Create a single port ram, called “dMem” (our data memory), using the Xilinx Block Memory IP in the same manner as described in Lab 4 Part 2. Set the width to 16 bits and the depth to 32768 (a total of 64 KiB). [Notice how this memory is not large enough to have a 16 bit address space since we are addressing by word rather than byte. Our processor has a “virtual” address space of 16 bits for the data memory but our physical system only has a 15 bit data address space due to memory limitations. In our top level design, we will ignore the top bit of our address when addressing this memory]. Disable any output registering and set the enables to always enabled. In the “Other Options” tab, set it to initialize using the data coe file generated in Part 1.

4.3 Debug Hints

- You can simulate your entire HDL wrapper and watch your entire system execute cycle by cycle. Very useful for testing your design. It will behave exactly the same way in the physical system. Be careful of sensitivity lists for processes to avoid mismatch vs. hardware.
- If Vivado decides to simulate post synthesis and dissolve your busses into wires, copy the top level HDL wrapper and the other HDL files into a new project. Generate the dMem and irMem using the IP Library like in lab 4. Modify the component names in the HDL wrapper. Now you can simulate the RTL version and add signals such as your state to get better waveforms that are easier to read as well as a properly enumerated state instead of one-hot encoded.

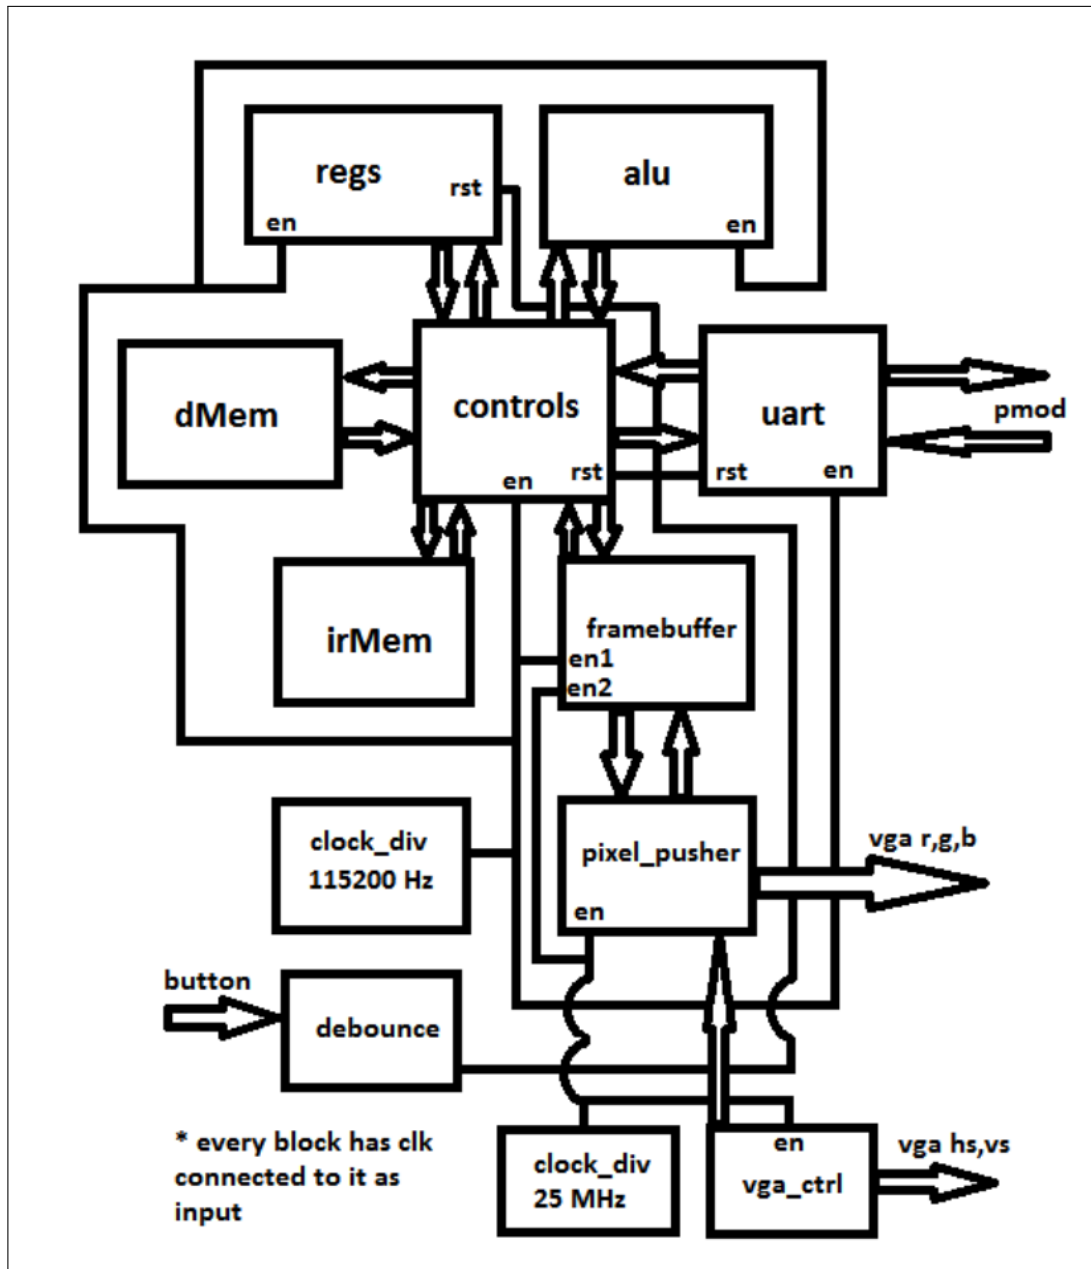


Figure 5.2: Lab 5 Simplified Top Level Block Diagram

5 Extra Credit

- (0.5 points) Implement a paging system for your instruction and data memories. Use the top unused bits for their addressing to mux IO to multiple chips to allow for in-place memory expansion of the system (think adding RAM to a computer).
- (0.5 points) Unlink the UART from all the other components' enable domain and make it have its own dedicated clock divider.
- (0.5 points) Make the max value for the UART clock divider come from \$r31 such that you can write

a value to \$r31 using the processor to change the UART baud rate.

- (1 point) Design and specify a video DMA block that takes as input a source base address in data memory, a destination base address in video memory, and a number of bytes. It then copies the memory from the source to the destination. You will need to convert your Data Memory to a dual port ram and add an instruction to the processor ISA by adding some states to the control FSM and modifying the assembler.
- (1 point) Implement the video DMA block that you have designed and test it with a sample program of your design.

6 GRISC Information

Table 5.3: GRISC ISA

Instruction	Opcode	Format	Meaning
add	0b00000	[op][reg1][reg2][reg3]	reg1 = reg2 + reg3
sub	0b00001	[op][reg1][reg2][reg3]	reg1 = reg2 - reg3
sll	0b00010	[op][reg1][reg2][*reg3]	reg1 = reg2 << 1
srl	0b00011	[op][reg1][reg2][*reg3]	reg1 = reg2 << 1
sra	0b00100	[op][reg1][reg2][*reg3]	reg1 = reg2 >>> 1
and	0b00101	[op][reg1][reg2][reg3]	reg1 = reg2 and reg3
or	0b00110	[op][reg1][reg2][reg3]	reg1 = reg2 or reg3
xor	0b00111	[op][reg1][reg2][reg3]	reg1 = reg2 xor reg3
slt	0b01000	[op][reg1][reg2][reg3]	reg1 = (reg2 < reg3) ? 1 : 0
sgt	0b01001	[op][reg1][reg2][reg3]	reg1 = (reg2 > reg3) ? 1 : 0
seq	0b01010	[op][reg1][reg2][reg3]	reg1 = (reg2 == reg3) ? 1 : 0
send (asip)	0b01011	[op][reg1][*reg2][*reg3]	sendUART(reg1[7:0])
recv (asip)	0b01100	[op][reg1][*reg2][*reg3]	reg1 = 0x00 & recvUART()
jr	0b01101	[op][reg1][*reg2][*reg3]	pc = reg1
wpix (asip)	0b01110	[op][reg1][reg2][*reg3]	framebuffer[reg1[11:0]] = reg2[15:0]
rpix (asip)	0b01111	[op][reg1][reg2][*reg3]	reg1 = framebuffer[reg2[11:0]]
beq	0b10000	[op][reg1][reg2][imm]	if(reg1 == reg2) pc = imm
bne	0b10001	[op][reg1][reg2][imm]	if(reg1 != reg2) pc = imm
ori	0b10010	[op][reg1][reg2][imm]	reg1 = reg2 or imm
lw	0b10011	[op][reg1][reg2][imm]	reg1 = dmem[reg2 + imm]
sw	0b10100	[op][reg1][reg2][imm]	dmem[reg2 + imm] = reg1
j	0b11000	[op][imm]	pc = imm
jal	0b11001	[op][imm]	ra = pc, pc = imm
clrscr (asip)	0b11010	[op][imm]	framebuffer[all] = 0
<i>Note: * means operand ignored</i>			

Table 5.4: GRISC Register File Organization

Register Index	Alias	Size (bits)	Special Behavior
0	\$zero	16	Resets to 0 every clock tick
1	\$pc	16	Program Counter
2	\$ra	16	Return Address for JAL
3-31	\$r3-\$r31	16	

Table 5.5: GRISC Instruction Decoding Table

Instruction Type	Opcode (4 downto 3)	Format
R	0b00 or 0b01	[op][reg1][reg2][reg3]
I	0b10	[op][reg1][reg2][imm]
J	0b11	[op][imm]

7 Credits

Sources

Acknowledgments

Gregory Leonberg

Original Author

Fall 2017

Stephen DiNicolantonio

Updates/Modification and L^AT_EX Design

Fall 2018